

# Machine Learning Research Paper 9

This study examines the interpretability of deep neural networks through gradient-based attribution methods. We develop a unified framework for explaining model predictions in computer vision tasks.