

Machine Learning Research Paper 13

An exploration of knowledge distillation techniques for model compression. We show that intermediate layer matching enables smaller student models to retain 95% of teacher accuracy.